

# NLP\_lab-7\_205229118\_Mahalakshmi.S

April 27, 2021

## 0.0.1 Lab7. Sentiment Analysis on Movie Reviews

In this lab, you will build Multinomial Naïve Bayes model for movie reviews from Rotton Tomotto Dataset.

### 0.1 EXERCISE-1

#### 0.1.1 1. Open the file, 'rotten\_tomato\_train.tsv' and read into a DataFrame

```
[75]: import pandas as pd
```

```
[76]: rotten_tomato_train = pd.read_csv('rotten_tomato_train.tsv', sep='\t')
```

#### 0.1.2 2. Print the basic statistics such as head, shape, describe, and columns

```
[77]: rotten_tomato_train.head()
```

```
[77]:
```

	PhraseId	SentenceId	Phrase \
0	1	1	A series of escapades demonstrating the adage ...
1	2	1	A series of escapades demonstrating the adage ...
2	3	1	A series
3	4	1	A
4	5	1	series

```
Sentiment
```

0	1
1	2
2	2
3	2
4	2

```
[78]: rotten_tomato_train.shape
```

```
[78]: (156060, 4)
```

```
[79]: rotten_tomato_train.describe
```

```
[79]: <bound method NDFrame.describe of
```

	PhraseId	SentenceId	\
0	1	1	

1	2	1
2	3	1
3	4	1
4	5	1
...	...	...
156055	156056	8544
156056	156057	8544
156057	156058	8544
156058	156059	8544
156059	156060	8544

	Phrase	Sentiment
0	A series of escapades demonstrating the adage ...	1
1	A series of escapades demonstrating the adage ...	2
2	A series	2
3	A	2
4	series	2
...	...	...
156055	Hearst 's	2
156056	forced avuncular chortles	1
156057	avuncular chortles	3
156058	avuncular	2
156059	chortles	2

[156060 rows x 4 columns]>

```
[80]: rotten_tomato_train.columns
```

```
[80]: Index(['PhraseId', 'SentenceId', 'Phrase', 'Sentiment'], dtype='object')
```

### 0.1.3 3. How many reviews exist for each sentiment?

```
[81]: review=rotten_tomato_train.groupby('Sentiment').count()
review.Phrase
```

```
[81]: Sentiment
0      7072
1     27273
2     79582
3     32927
4      9206
Name: Phrase, dtype: int64
```

## 0.2 EXERCISE-2

0.2.1 1. Extract 200 reviews for each sentiment, store them into a new dataframe and create a smaller dataset. Save this dataframe in a new file, say, “small\_rotten\_train.csv”.

```
[82]: a=rotten_tomato_train.loc[rotten_tomato_train.Sentiment == 0]
      b=rotten_tomato_train.loc[rotten_tomato_train.Sentiment == 1]
      c=rotten_tomato_train.loc[rotten_tomato_train.Sentiment == 2]
      d=rotten_tomato_train.loc[rotten_tomato_train.Sentiment == 3]
      e=rotten_tomato_train.loc[rotten_tomato_train.Sentiment == 4]
```

```
[83]: small_rotten_train=pd.concat([a[:200],b[:200],c[:200],d[:200],e[:200]])
```

## 0.3 EXERCISE-3

0.3.1 1. Open the file, “small\_rotten\_train.csv”.

```
[84]: small_rotten_train
```

```
[84]:
```

	PhraseId	SentenceId	Phrase \
101	102	3	would have a hard time sitting through this one
103	104	3	have a hard time sitting through this one
157	158	5	Aggressive self-glorification and a manipulati...
159	160	5	self-glorification and a manipulative whitewash
201	202	7	Trouble Every Day is a plodding mess .
...	...	...	...
3744	3745	142	amazing slapstick
3745	3746	142	amazing
3847	3848	147	When cowering and begging at the feet a scruff...
3866	3867	147	gives her best performance since Abel Ferrara ...
3993	3994	151	Spielberg 's realization of a near-future Amer...

	Sentiment
101	0
103	0
157	0
159	0
201	0
...	...
3744	4
3745	4
3847	4
3866	4
3993	4

```
[1000 rows x 4 columns]
```

**0.3.2 2. The review text are stored in “Phrase” column. Extract that into a separate DataFrame, say “X”.**

```
[85]: X = small_rotten_train.Phrase
```

**0.3.3 3. The “sentiment” column is your target, say “y”.**

```
[86]: y = small_rotten_train.Sentiment
```

**0.3.4 4. Perform pre-processing: convert into lower case, remove stop words and lemmatize. The following function will help.**

```
[87]: import nltk
      from nltk.corpus import stopwords
      nltk.download('stopwords')
      stop_words = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Mahalakshmi\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
[88]: from nltk.stem import WordNetLemmatizer
      lemmatizer = WordNetLemmatizer()
```

```
[89]: def clean_review(review):
      tokens = review.lower().split()
      filtered_tokens = [lemmatizer.lemmatize(w) for w in tokens if w not in
      ↪stop_words]
      return " ".join(filtered_tokens)
```

**0.3.5 5. Apply the above function to X**

```
[90]: temp=X.tolist()
      fax=[]
```

```
[91]: for i in temp:
      fax.append(clean_review(i))
      n_X=pd.Series(fax)
```

**0.3.6 6. Split X and y for training and testing (Use 20% for testing)**

```
[92]: from sklearn.model_selection import train_test_split
```

```
[93]: X_train,X_test,y_train,y_test = train_test_split(n_X,y,train_size=0.
      ↪8,test_size=0.2)
```

**0.3.7 7. Create TfidfVectorizer as below and perform vectorization on X\_train using fit\_transform() method.**

```
[94]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[95]: tf=TfidfVectorizer(min_df=3, max_features=None,ngram_range=(1, 2), use_idf=1)
      tf
```

```
[95]: TfidfVectorizer(min_df=3, ngram_range=(1, 2), use_idf=1)
```

```
[96]: tfi=tf.fit_transform(X_train)
      tfi.shape
```

```
[96]: (800, 904)
```

**0.3.8 8. Create MultinomialNB model and perform training using X\_train\_lemmatized and y\_train.**

```
[97]: from sklearn.feature_extraction.text import CountVectorizer
      cv = CountVectorizer()
```

```
[98]: X_train_mnb = cv.fit_transform(X_train)
      X_test_mnb = cv.transform(X_test)
```

```
[99]: from sklearn.naive_bayes import MultinomialNB
```

```
[100]: clf = MultinomialNB()
```

```
[101]: clf.fit(X_train_mnb,y_train)
```

```
[101]: MultinomialNB()
```

**0.3.9 9. Perform validation on X\_test lemmatized and predict output**

```
[102]: y_lem_pred = clf.predict(X_test_mnb)
      y_lem_pred
```

```
[102]: array([2, 0, 4, 0, 3, 4, 1, 4, 3, 3, 4, 0, 0, 1, 4, 4, 1, 1, 2, 3, 3, 4,
        1, 3, 0, 2, 4, 3, 1, 3, 0, 1, 2, 3, 3, 4, 3, 3, 1, 1, 0, 2, 3, 3,
        1, 2, 1, 4, 3, 1, 4, 1, 3, 4, 1, 2, 3, 2, 2, 4, 4, 0, 3, 4, 2, 3,
        1, 2, 1, 4, 0, 1, 1, 1, 0, 0, 3, 0, 4, 0, 1, 3, 2, 1, 4, 3, 3, 0,
        2, 4, 0, 3, 3, 2, 3, 1, 4, 1, 2, 3, 2, 0, 4, 1, 1, 3, 3, 1, 1, 4,
        0, 2, 0, 2, 2, 0, 2, 0, 0, 2, 3, 3, 1, 0, 2, 0, 3, 1, 1, 1, 3, 3,
        4, 3, 4, 1, 2, 2, 1, 2, 3, 0, 1, 0, 0, 2, 0, 3, 1, 3, 4, 2, 2, 2,
        1, 4, 1, 4, 2, 4, 3, 2, 0, 3, 3, 2, 3, 1, 2, 1, 3, 0, 4, 2, 0, 1,
        4, 1, 2, 0, 2, 2, 4, 1, 3, 3, 1, 1, 3, 2, 4, 3, 3, 0, 3, 1, 1, 1,
        3, 3], dtype=int64)
```

### 0.3.10 10. Print classification\_report and accuracy score.

```
[103]: from sklearn.metrics import classification_report
```

```
[104]: print(classification_report(y_test,y_lem_pred))
```

	precision	recall	f1-score	support
0	0.88	0.68	0.77	41
1	0.62	0.66	0.64	44
2	0.58	0.52	0.55	42
3	0.47	0.69	0.56	35
4	0.81	0.68	0.74	38
accuracy			0.65	200
macro avg	0.67	0.65	0.65	200
weighted avg	0.67	0.65	0.65	200

```
[105]: from sklearn.metrics import accuracy_score
```

```
[106]: accuracy_score(y_test,y_lem_pred)
```

```
[106]: 0.645
```

### 0.3.11 EXERCISE-4

#### 0.3.12 1. Open, 'rotten\_tomato\_test.tsv' file into dataframe

```
[107]: rotten_tomato_test = pd.read_csv('rotten_tomato_test.tsv', sep='\t')
```

```
[108]: rotten_tomato_test.head()
```

```
[108]:
```

	PhraseId	SentenceId	Phrase
0	156061	8545	An intermittently pleasing but mostly routine ...
1	156062	8545	An intermittently pleasing but mostly routine ...
2	156063	8545	An
3	156064	8545	intermittently pleasing but mostly routine effort
4	156065	8545	intermittently pleasing but mostly routine

```
[109]: rotten_tomato_test.shape
```

```
[109]: (66292, 3)
```

#### 0.3.13 2. Clean this test data, using the function clean\_review(), as before.

```
[110]: X_c = rotten_tomato_test.Phrase
```

```
[111]: t_temp=X_c.tolist()
t_fax=[]
for i in t_temp:
    t_fax.append(clean_review(i))
cr_X=pd.Series(t_fax)
```

```
[112]: cr_X
```

```
[112]: 0      intermittently pleasing mostly routine effort .
1      intermittently pleasing mostly routine effort
2
3      intermittently pleasing mostly routine effort
4      intermittently pleasing mostly routine
...
66287      long-winded , predictable scenario .
66288      long-winded , predictable scenario
66289      long-winded ,
66290      long-winded
66291      predictable scenario
Length: 66292, dtype: object
```

### 0.3.14 3. Build TFIDF values using transform() method.

```
[113]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[114]: tf_1=TfidfVectorizer(use_idf=True,ngram_range=(1,3),min_df = 1)
tf_1
```

```
[114]: TfidfVectorizer(ngram_range=(1, 3))
```

```
[116]: vec=tf_1.transform(cr_X)
vec
```

```
[116]: <66292x904 sparse matrix of type '<class 'numpy.float64'>'
      with 68274 stored elements in Compressed Sparse Row format>
```